

# VISTA-MPC

Vision-Language Informed Socially Attentive  
Model Predictive Control for Robot Navigation

Jiaxi Huo

University of Technology Sydney

January 29, 2026

# Outline

- 1 Motivation
- 2 System Architecture
- 3 SARL: Socially Attentive Reinforcement Learning
- 4 MPC Formulation
- 5 VLM Integration
- 6 Three Synergy Mechanisms
- 7 ROS 2 Implementation
- 8 Configuration & Deployment
- 9 Algorithm Summary
- 10 Discussion & Conclusion

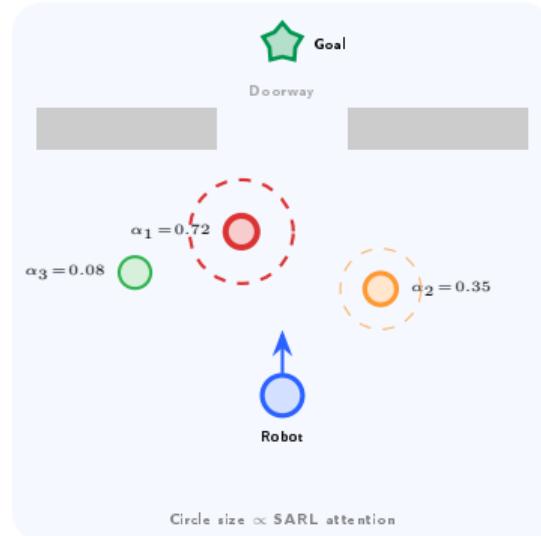
# The Challenge of Social Robot Navigation

**Problem:** Robots navigating in human-populated environments must respect *implicit social norms*, not just avoid collisions.

## Key challenges:

- Variable crowd density and scene structure
- Not all pedestrians are equally important
- Context-dependent social behavior ( $\text{doorway} \neq \text{open space}$ )
- Real-time control under uncertainty
- Need for *interpretable* decisions

**Our Insight:** Combine learned social attention (SARL), scene understanding (VLM), and optimal control (MPC) in a principled hierarchy.



# Limitations of Existing Approaches

## Reactive Methods (DWA, RVO2)

- Treat all pedestrians equally
- No scene context understanding
- Myopic — no lookahead

## End-to-End RL

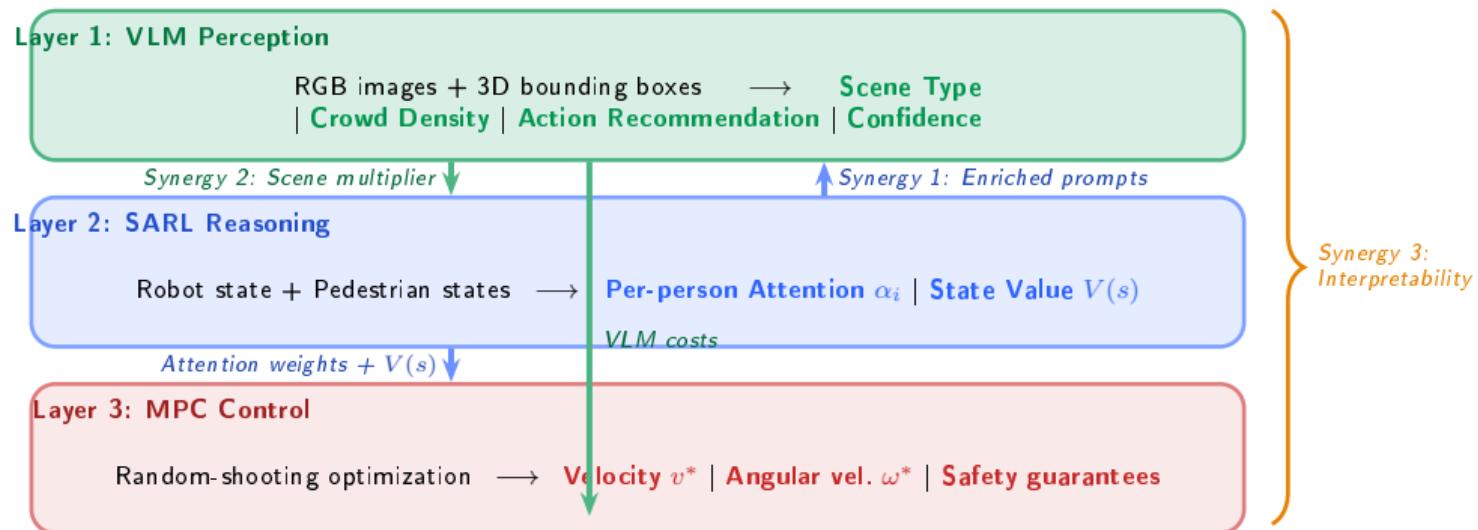
- Black-box decisions
- Sim-to-real transfer gap
- No safety guarantees

## VISTA-MPC + SARNL (Ours)

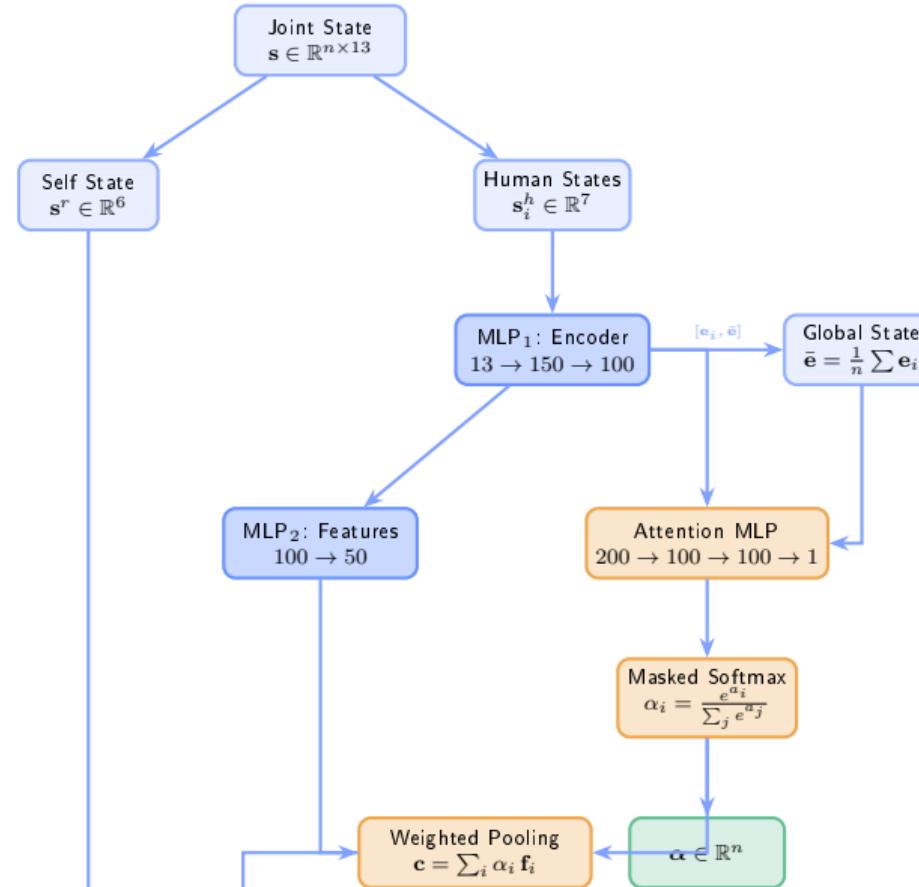
- **Attention-weighted** social reasoning
- **Scene-adaptive** via VLM
- **Predictive** MPC horizon (3 s)
- **Interpretable** explanations
- **Safety guarantees** via hard constraints
- **Graceful fallback** when data stale

**Key Idea:** Use RL for *what to attend to*, VLM for *scene understanding*, and MPC for *safe optimal control*.

# Three-Layer Hierarchical Architecture



# SARL Value Network Architecture



# SARL: State Representation & Coordinate Transform

**Joint State Vector** (per human  $i$ ):

$$\mathbf{s}_i = \left[ \underbrace{d_g, v_{\text{pref}}, \theta, r}_{\text{self state}}, \underbrace{v_x^r, v_y^r}_{\text{velocity}}, \underbrace{p_x^i, p_y^i, v_x^i, v_y^i, r^i, d_a, r_\Sigma}_{\text{human state}} \right]$$

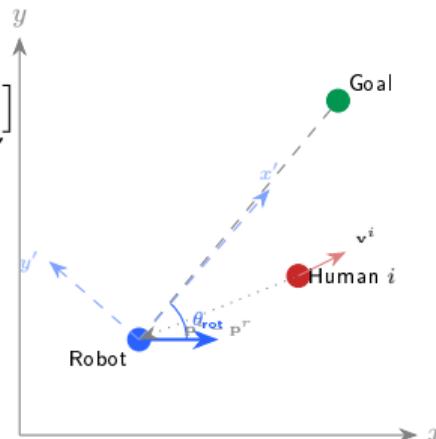
**Goal-Oriented Rotation:**

$$\theta_{\text{rot}} = \text{atan2}(g_y - p_y^r, g_x - p_x^r)$$

$$\mathbf{R} = \begin{bmatrix} \cos \theta_{\text{rot}} & \sin \theta_{\text{rot}} \\ -\sin \theta_{\text{rot}} & \cos \theta_{\text{rot}} \end{bmatrix}$$

Applied to all velocity and position vectors:

$$\mathbf{v}_{\text{rot}} = \mathbf{R} \mathbf{v}_{\text{world}}, \quad \mathbf{p}_{\text{rot}}^i = \mathbf{R} (\mathbf{p}_{\text{world}}^i - \mathbf{p}^r)$$



Why Rotate?

Goal-centric frame makes the policy *translation and rotation invariant* — the same network weights generalize across positions.

# Self-Attention Mechanism

**Stage 1:** Encode each human–robot pair:

$$\mathbf{e}_i = \text{MLP}_1(\mathbf{s}_i) \in \mathbb{R}^{100}$$

**Stage 2:** Compute global context:

$$\bar{\mathbf{e}} = \frac{1}{n} \sum_{i=1}^n \mathbf{e}_i$$

**Stage 3:** Attention scores with global state:

$$a_i = \text{MLP}_{\text{attn}}([\mathbf{e}_i, \bar{\mathbf{e}}]) \in \mathbb{R}$$

**Stage 4:** Masked softmax normalization:

$$\alpha_i = \frac{\exp(a_i) \cdot \mathbb{1}[\mathbf{s}_i^h \neq \mathbf{0}]}{\sum_{j=1}^n \exp(a_j) \cdot \mathbb{1}[\mathbf{s}_j^h \neq \mathbf{0}]}$$

**Stage 5:** Weighted feature aggregation:

## Interpretation

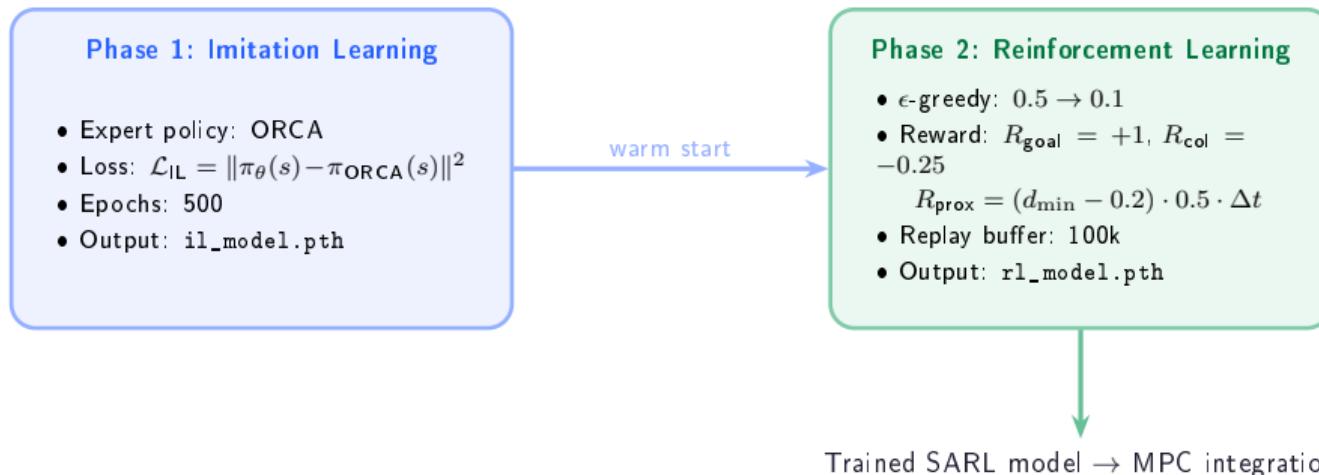
- $\alpha_i$  captures *behavioral importance* of pedestrian  $i$
- Not just proximity — considers velocity, heading, relative position
- Zero-masking handles variable crowd sizes
- Global state  $\bar{\mathbf{e}}$  provides crowd context to each attention score

## Value Function

$$V(s) = \text{MLP}_3([\mathbf{s}^r, \mathbf{c}]) \in \mathbb{R}$$

- Encodes long-term navigation quality
- Human–MPC + visual input

# Two-Phase Training Pipeline



Simulation: Custom Gym environment with RVO2 humans.  
Circle-crossing scenario, configurable crowd size (5–20 humans).  
Discount factor  $\gamma = 0.99$ , batch size 64, target network updated every 500 steps.

# Model Predictive Control: Problem Formulation

## Unicycle Dynamics

$$\begin{aligned}x_{k+1} &= x_k + v_k \cos(\theta_k) \cdot \Delta t \\y_{k+1} &= y_k + v_k \sin(\theta_k) \cdot \Delta t \\\theta_{k+1} &= \theta_k + \omega_k \cdot \Delta t\end{aligned}$$

## MPC Optimization

$$(v^*, \omega^*) = \arg \min_{v_{0:N-1}, \omega_{0:N-1}} \sum_{k=0}^{N-1} \underbrace{J_k^{\text{stage}}}_{\text{stage cost}} + \underbrace{J_N^{\text{term}}}_{\text{terminal cost}}$$

## Subject to:

- $0 \leq v_k \leq v_{\max}$
- $|\omega_k| \leq \omega_{\max}$
- $d_{\text{obs}}(\mathbf{x}_k) \geq d_{\text{hard}}$  (safety)
- Unicycle dynamics (above)

$N = 15$  steps,  $\Delta t = 0.2$  s (horizon = 3.0 s), 100 random rollouts, 10 Hz control rate

# MPC Cost Function: Full Breakdown

$$J_k^{\text{stage}} = \underbrace{J_k^{\text{goal}}}_{\text{Goal}} + \underbrace{J_k^{\text{social}}}_{\text{SARL Social}} + \underbrace{J_k^{\text{obs}}}_{\text{Obstacle}} + \underbrace{J_k^{\text{smooth}}}_{\text{Smoothness}} + \underbrace{J_k^{\text{VLM}}}_{\text{VLM}}$$

## Goal Cost:

$$J_k^{\text{goal}} = w_g \| (x_k - g_x, y_k - g_y) \|^2$$

## SARL Attention-Weighted Social Cost:

$$J_k^{\text{social}} = \sum_{i=1}^n \frac{w_s \cdot n \cdot \alpha_i \cdot m_{\text{attn}}}{d(\mathbf{x}_k, \mathbf{p}_i) + \epsilon}$$

where  $\alpha_i$  = SARL attention,  $m_{\text{attn}}$  = scene multiplier

## Smoothness Cost:

$$J_k^{\text{smooth}} = w_{\text{sm}} [(v_k - v_{k-1})^2 + (\omega_k - \omega_{k-1})^2]$$

## Obstacle Cost:

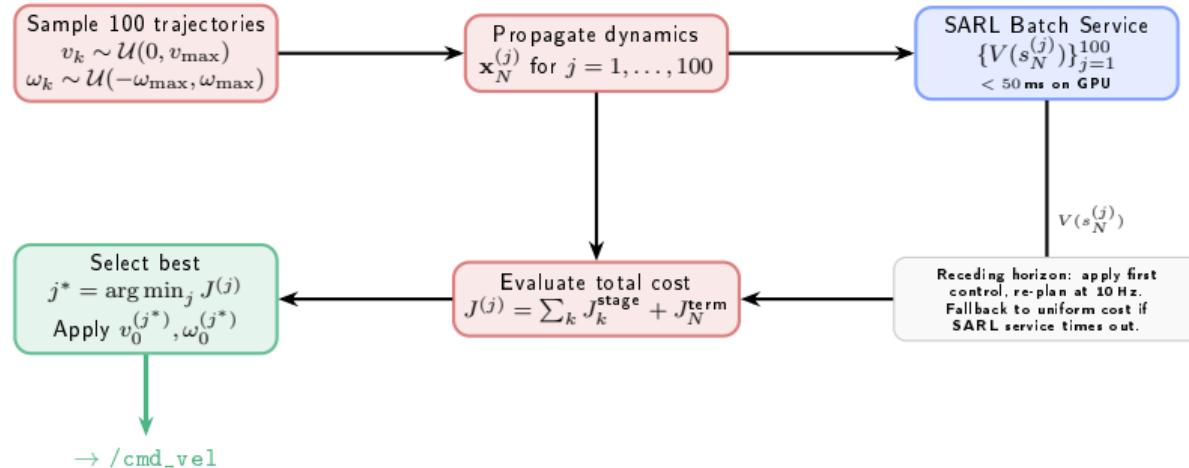
$$J_k^{\text{obs}} = \begin{cases} +\infty & d < d_{\text{hard}} \\ 1000 & d_{\text{hard}} \leq d < d_{\text{min}} \\ \frac{w_o}{d + 0.1} & d \geq d_{\text{min}} \end{cases}$$

## SARL Terminal Cost:

$$J_N^{\text{term}} = -w_t \cdot m_{\text{term}} \cdot V(s_N)$$

Negated: higher  $V(s) \Rightarrow$  lower cost  
 $\Rightarrow$  MPC seeks high-value terminal states

# Random-Shooting MPC with SARNL Batch Evaluation



# Vision-Language Model: Scene Understanding

## VLM Input (Prompt Assembly):

- RGB camera image
- 3D bounding boxes of detected pedestrians
- Robot position and goal
- **SARL attention weights** per person (Synergy 1)

## VLM Output (Structured):

- **Scene type**: doorway, corridor, crossing, queue, open\_space
- **Crowd density**: empty, sparse, medium, dense
- **Recommended action**: go\_ahead, slow\_down, yield, wait
- **Confidence**: [0, 1]
- **Min. personal distance**: adaptive threshold

## Trigger Conditions:

- None

### Example VLM Prompt (with SARL)

Robot at (2.1, 3.4), goal (8.0, 1.0)

Pedestrians (3):

1. ID=001, Pos=(3.5, 4.1)  
Vel=(0.2, -0.1) m/s  
**SARL\_ATTN=0.72**  
**RL\_RISK=HIGH**
2. ID=002, Pos=(5.2, 2.8)  
Vel=(0.0, 0.3) m/s  
**SARL\_ATTN=0.20**  
**RL\_RISK=LOW**
3. ID=003, Pos=(1.0, 5.5)  
Vel=(-0.1, 0.0) m/s  
**SARL\_ATTN=0.08**  
**RL\_RISK=LOW**

# VLM Additional Cost Terms for MPC

The VLM translator converts semantic outputs into four MPC-compatible cost terms:

## 1. Directional Cost ( $w_{\text{vlm-dir}}$ ):

- Penalizes heading deviation from VLM-recommended direction
- E.g., “pass on the left”

## 2. Action Cost ( $w_{\text{vlm-act}}$ ):

- Penalizes velocity when VLM says “stop and wait” or “yield”

## 3. Scene Cost ( $w_{\text{vlm-scene}}$ ):

- Scene-specific penalties (doorway narrowness, crossing density)

## 4. Personal Space Cost ( $w_{\text{vlm-pers}}$ ):

- Distance violation when closer than VLM-recommended minimum

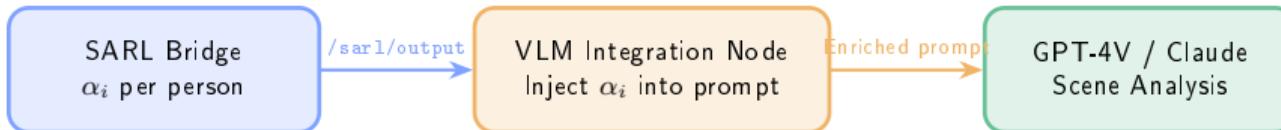
## VLM Cost in MPC

$$\begin{aligned} J_k^{\text{vlm}} = & w_{\text{dir}} \cdot C_{\text{dir}}(\theta_k) \\ & + w_{\text{act}} \cdot C_{\text{act}}(v_k) \\ & + w_{\text{scene}} \cdot C_{\text{scene}}(\mathbf{x}_k) \\ & + w_{\text{pers}} \cdot C_{\text{pers}}(d_{\min}^k) \end{aligned}$$

## Enable/Disable VLM

VLM can be toggled via launch argument `enable_vlm`. When disabled, MPC uses only SARL attention + standard costs. System remains fully functional.

# Synergy 1: SARNL Attention Enriches VLM Prompts



**Without SARNL enrichment:** "Person\_001: distance=2.5m, motion=walking" → VLM treats all equally by distance

**With SARNL enrichment:** "Person\_001: distance=2.5m, **SARNL\_ATTENTION=0.72, RL\_RISK=HIGH**" → VLM understands *behavioral* importance

**Result:** VLM recommendations are informed by what the RL policy has *learned* matters, not just geometric proximity.

## Synergy 2: VLM Scene Context Conditions SARNL Weights

VLM provides **scene\_type** and **crowd\_density**.

These modulate how much MPC trusts SARNL via `getSceneMultiplier()`:

$$J_k^{\text{social}} = \sum_i \frac{w_s \cdot n \cdot \alpha_i \cdot m_{\text{attn}}}{d_i + \epsilon}$$

$$J_N^{\text{term}} = -w_t \cdot m_{\text{term}} \cdot V(s_N)$$

### Intuition:

- **Doorway** (constrained): Amplify SARNL → trust learned social behavior more
- **Open space**: Reduce SARNL → simple distance-based cost suffices

**Scene Multiplier Matrix**

Scene	$m_{\text{attn}}$	$m_{\text{term}}$
Doorway (dense)	2.5	1.2
Corridor (dense)	2.0	1.0
Corridor (medium)	1.5	0.7
Crossing (medium)	1.5	0.7
Crossing (sparse)	1.0	0.5
Open space (sparse)	0.5	0.3



# Synergy 3: Combined Interpretability

The NavigationExplanation message merges VLM and SARNL outputs into a *single unified explanation*:

## Contents:

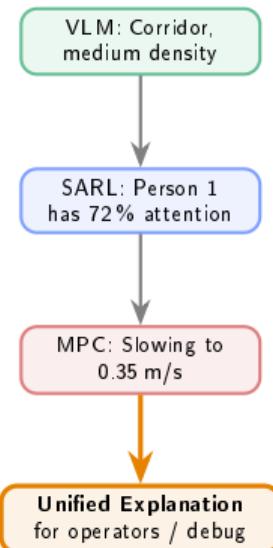
- VLM: scene type, crowd density, recommended action, confidence
- SARNL: per-person attention weights, distances, motion labels
- Control: applied  $v_{\max}$ ,  $w_{\text{social}}$ , scene multiplier
- **Human-readable narrative**

**Published to:** /navigation/explanation

## Example Explanation

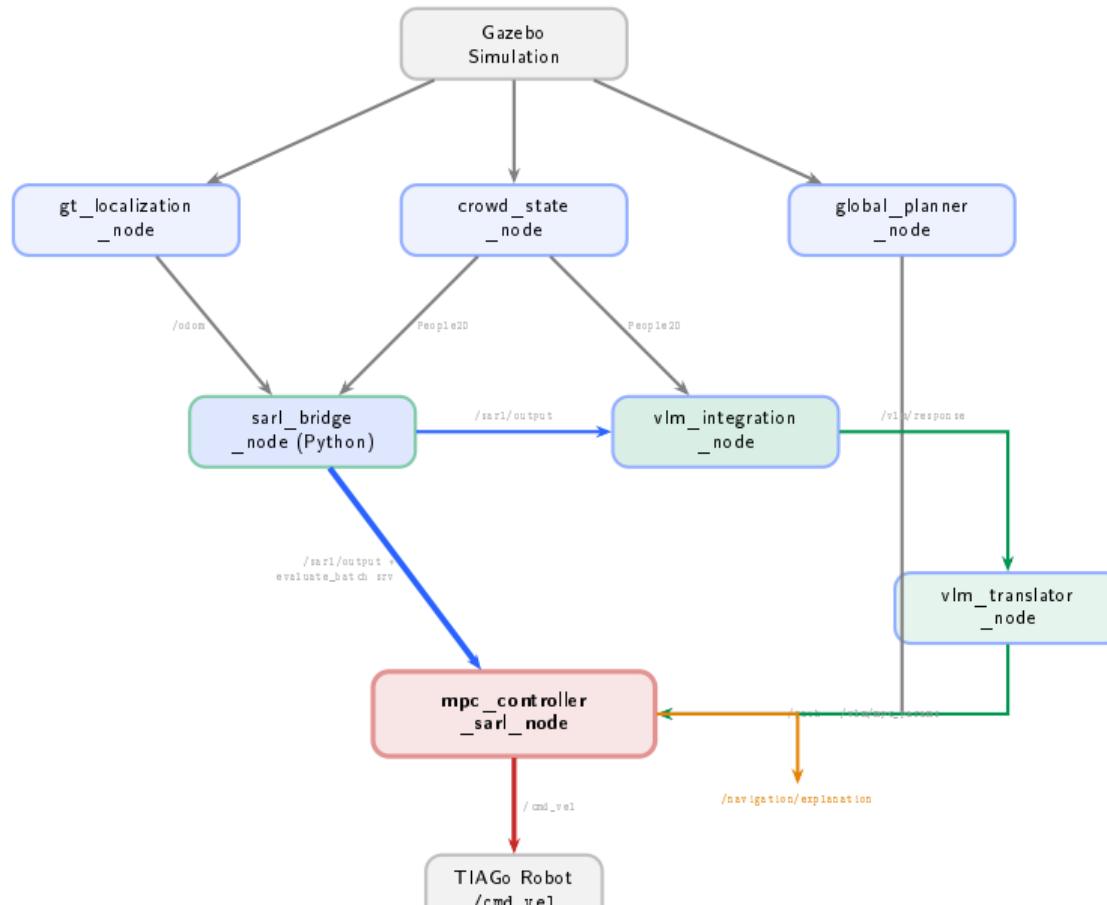
```
Scene: corridor (medium) | SARNL active:  
V(s)=0.35, attention_mult=1.5  
Focus: person_001 (w=0.72, HIGH)  
Action: slow_down | v_max=0.35 m/s
```

## Why Interpretability Matters



Enables trust calibration, debugging, and human oversight of autonomous navigation decisions.

# ROS2 Node Architecture



# SARL Bridge Node: Real-Time Inference

## Dual Interface Design:

### Interface 1: Topic Publisher (10 Hz)

- Subscribes: `/odom`, `/person_info`
- Publishes: `/sarl/output`
- Content:  $\{\alpha_i\}$ ,  $V(s)$ , person IDs
- RViz markers: colored spheres  
(green → yellow → red  $\propto$  attention)

### Interface 2: Batch Service

- Service: `/sarl/evaluate_batch`
- Input: 100 terminal states from MPC
- Output:  $\{V(s_N^{(j)})\}_{j=1}^{100}$
- Latency: < 50 ms with PyTorch

## Implementation Details:

- Loads `rl_model.pth` from CrowdNav
- Replicates ValueNetwork architecture exactly
- Goal-oriented `rotate()` coordinate transform
- Unicycle  $\rightarrow$  holonomic velocity conversion:  
 $v_x = v \cos(\theta), \quad v_y = v \sin(\theta)$

## Staleness Handling:

- SARL data timestamped at publication
- MPC checks: if  $\text{age} > 0.5 \text{ s} \Rightarrow$  fallback to uniform social cost
- Ensures safety even if SARL node crashes

## RViz Attention Markers:

- Sphere radius:  $0.2 + \alpha_i \times 0.6 \text{ m}$

# Social Contract Helper: Rule-Based Safety Layer

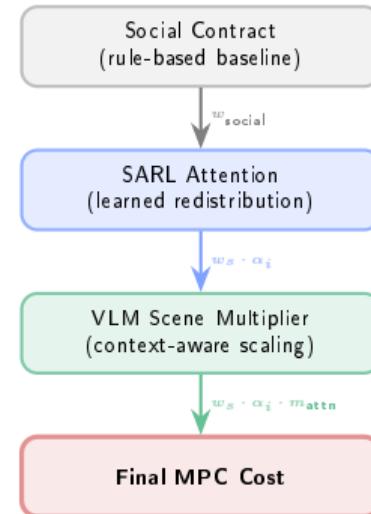
**Purpose:** Proximity-based heuristic that adjusts MPC parameters *before* optimization, providing an additional safety layer on top of SARL attention.

## Adjustable Parameters:

- $v_{\max}$ : reduced when people nearby
- $w_{\text{social}}$ : increased when person in path
- $d_{\min}^{\text{person}}$ : minimum distance threshold
- `person_in_front`: boolean flag for blocking detection

## Integration with SARL:

- Social contract sets *baseline*  $w_{\text{social}}$
- SARL attention *redistributes* this weight across pedestrians
- VLM scene multiplier *scales* the combined result



# Key Configuration Parameters

## MPC Parameters (`mpc_sarl.yaml`)

Parameter	Value	Unit
$\Delta t$	0.2	s
$N$ (horizon)	15	steps
Rollouts	100	—
Control rate	10	Hz
$w_{goal}$	2.0	—
$w_{social}$	1.0	—
$w_{obstacle}$	3.0	—
$w_{smooth}$	0.1	—
$w_{sarl-attn}$	1.0	—
$w_{sarl-term}$	0.5	—
Staleness timeout	0.5	s
$d_{hard}$	0.2	m
$d_{min}^{obs}$	0.3	m

## SARL Bridge (`sarl_bridge.yaml`)

Parameter	Value
Model path	<code>rl_model.pth</code>
MLP <sub>1</sub> dims	[150, 100]
MLP <sub>2</sub> dims	[100, 50]
MLP <sub>3</sub> dims	[150,100,100,1]
Attn dims	[100, 100, 1]
Global state	True
Publish rate	10 Hz

## Launch Command

```
ros2 launch social_mpc_nav  
  mpc_sarl_vlm.launch.py  
  enable_vlm:=true  
  log_mpc_to_csv:=true  
  goal_x:=10.0 goal_y:=5.0
```

# VISTA-MPC + SARN: Complete Algorithm

**Require:** Trained SARN model, VLM API, MPC config

```
1: Initialize ROS 2 nodes, load SARN weights
2: while navigation active (at 10 Hz) do
3:   Get robot state ( $x, y, \theta, v$ ) from odometry
4:   Get pedestrian states  $\{\mathbf{p}_i, \mathbf{v}_i\}$  from tracker
5:   // SARN Inference (10 Hz)
6:   Rotate states to goal-centric frame
7:   Compute  $\{\alpha_i\}, V(s)$  via ValueNetwork
8:   Publish /sarn/output
9:   // VLM (event-triggered)
10:  if trigger condition met then
11:    Assemble prompt with SARN  $\alpha_i$ 
12:    Query VLM  $\rightarrow$  scene type, density, action
13:  end if
14:  // MPC Optimization
15:  Get scene multiplier  $m_{\text{attn}}, m_{\text{term}}$ 
16:  Sample 100 random trajectories
17:  Call evaluate_batch  $\rightarrow \{V(s_N^{(j)})\}$ 
18:  Evaluate  $J^{(j)}$  with all cost terms
19:   $j^* \leftarrow \arg \min_j J^{(j)}$ 
20:  Apply  $v_0^{(j*)}, \omega_0^{(j*)}$  to /cmd_vel
21:  Publish explanation
```

Syn. 1

Syn. 2

Syn. 3

## Key Properties:

- **Anytime:** Best trajectory available at any point
- **Graceful degradation:** VLM timeout  $\rightarrow$  SARN-only; SARN stale  $\rightarrow$  uniform cost
- **Safety:** Hard distance constraints always enforced
- **Interpretable:** Every decision is explainable

## Computation Budget

SARN inference:  $\sim 5$  ms  
Batch evaluation:  $< 50$  ms  
MPC optimization:  $\sim 10$  ms  
VLM query: 1–3 s (async)

Total:  $< 100$  ms / cycle

# Key Contributions

## ① Hierarchical VLM–SARL–MPC integration

First framework combining learned social attention, vision-language scene understanding, and model predictive control in a principled three-layer hierarchy.

## ② Three synergy mechanisms

- S1** SARL attention enriches VLM prompts for better scene analysis
- S2** VLM scene context adaptively scales SARL cost contributions
- S3** Combined interpretability for human-readable navigation explanations

## ③ SARL terminal value for MPC

Learned value function  $V(s_N)$  provides long-horizon social reasoning beyond MPC's finite horizon (3 s), evaluated in batch for 100 rollouts.

## ④ Robust real-time system

ROS 2 implementation with graceful fallback (SARL stale  $\rightarrow$  uniform cost), hard safety constraints, and full CSV logging for experiment analysis.

# Advantages of the Hybrid Approach

Property	Pure RL	Pure MPC	VISTA-MPC+SARL
Social attention	✓	✗	✓
Scene understanding	✗	✗	✓ (VLM)
Safety guarantees	✗	✓	✓
Interpretability	✗	~	✓
Long-horizon reasoning	✓ ( $V(s)$ )	~ (finite $N$ )	✓ (both)
Scene adaptation	Fixed policy	Manual tuning	Automatic
Graceful fallback	✗	✓	✓

VISTA-MPC + SRL combines the **social intelligence** of deep RL, the **semantic reasoning** of VLMs, and the **safety & optimality** of MPC.

# Future Work

## Short-Term

- Real robot deployment on TIAGo platform
- Systematic ablation study:  
VLM-only vs SARNL-only vs combined
- Human-subject evaluation of social compliance
- Benchmark on standard SocNav datasets

## Long-Term

- End-to-end fine-tuning of SARNL with VLM feedback
- Pedestrian trajectory prediction integration
- Multi-robot coordination with shared attention
- CasADi-based gradient MPC replacing random shooting
- Smaller on-device VLMs for edge deployment

**Thank You!**

Questions?

Code: [github.com/huojiaxi/RL-Social-Nav](https://github.com/huojiaxi/RL-Social-Nav)

# Appendix: Message Definitions

## SARLOutput.msg

```
time stamp
float32[] attention_weights
string[] person_names
float32 state_value
float32 robot_x
float32 robot_y
float32 robot_yaw
float32 inference_time_ms
bool is_valid
```

## EvaluateSARLBatch.srv

```
# Request
float32[] robot_states
float32 goal_x
float32 goal_y
float32[] people_states
int32 num_rollouts
int32 num_people
--
# Response
float32[] values
float32 inference_time_ms
bool success
```

## NavigationExplanation

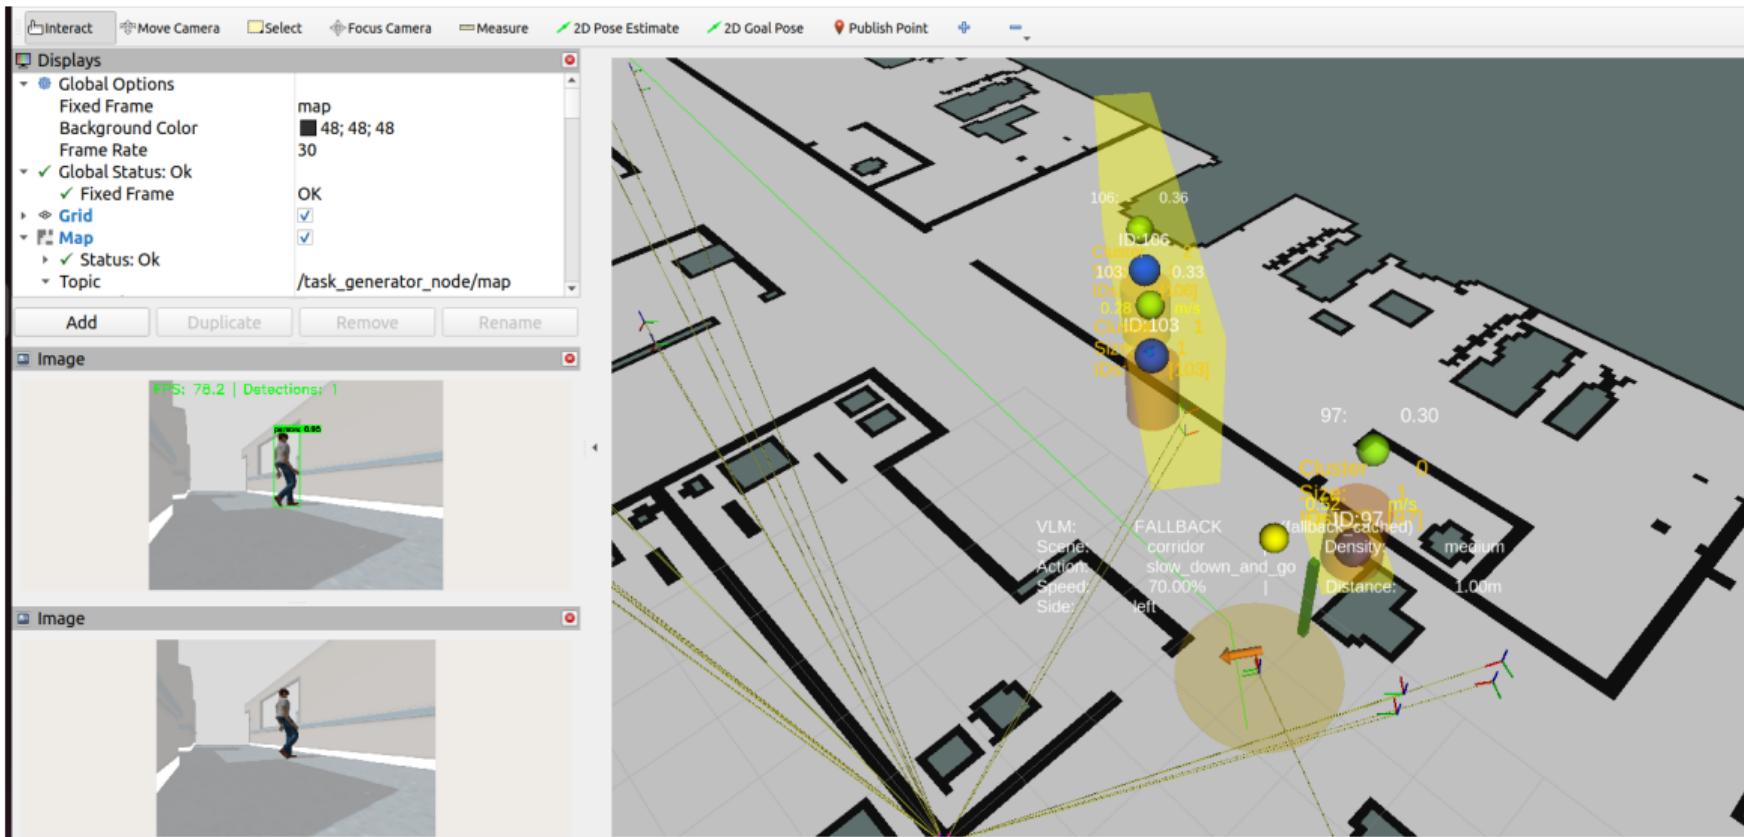
```
string scene_type
string crowd_density
string recommended_action
float32 vlm_confidence
string[] person_names
float32[] attention_weights
float32[] person_distances
string[] person_motions
float32 state_value
float32 applied_v_max
float32 applied_w_social
float32 sarl_scene_mult
string explanation_text
```

# Appendix: SARNL Network Dimensions

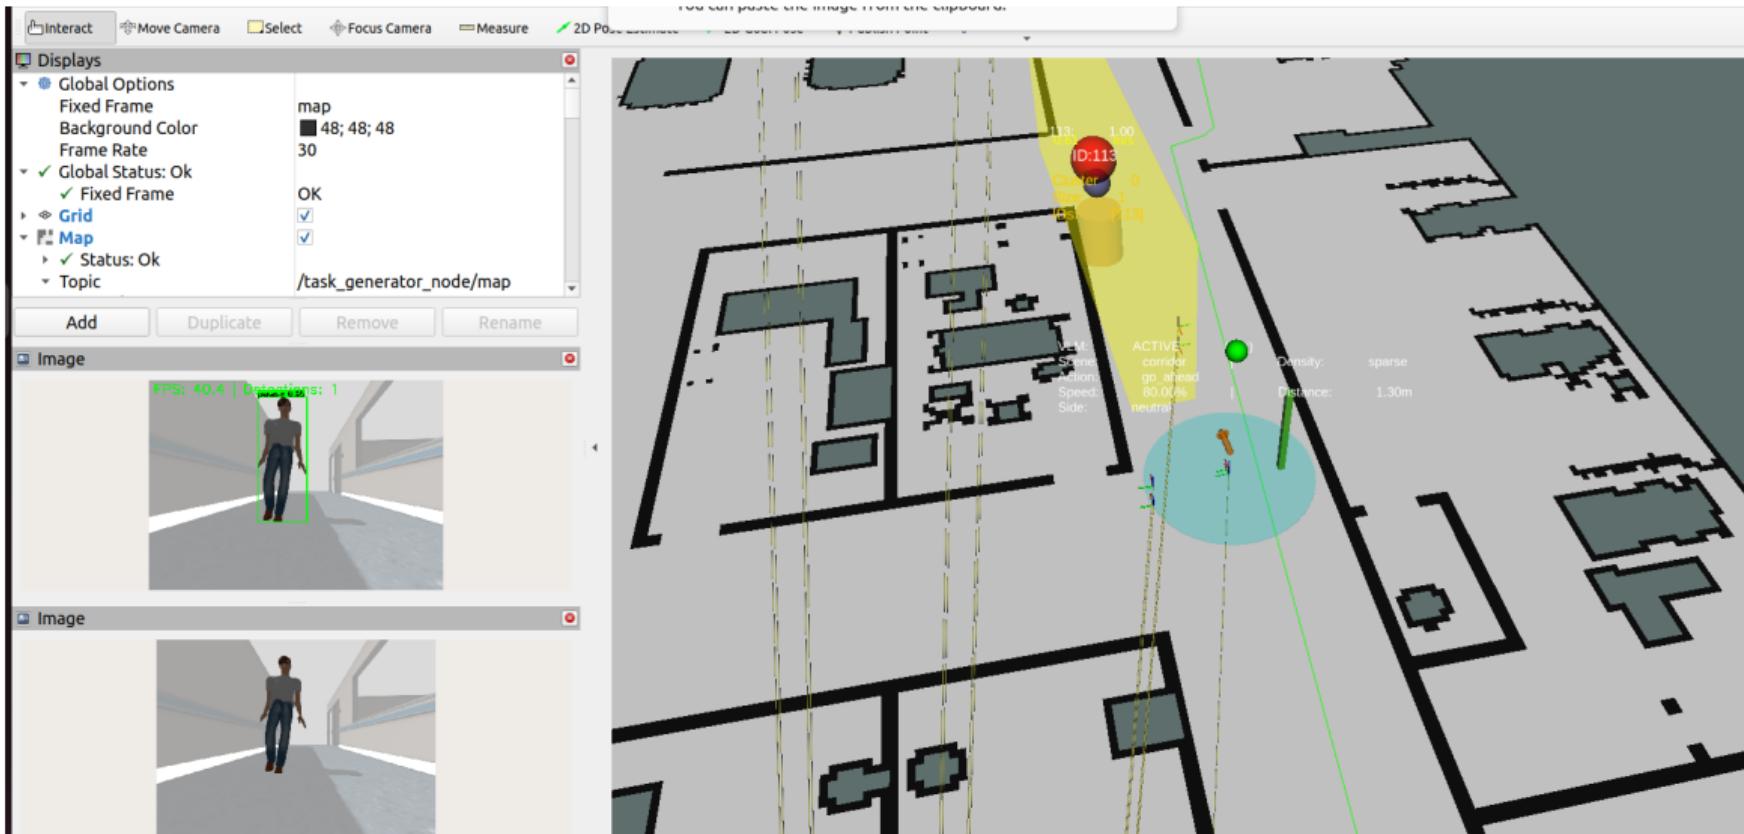
Component	Layer	Input	Output	Activation
MLP <sub>1</sub> (Encoder)	Linear + ReLU	13	150	ReLU
	Linear + ReLU	150	100	ReLU
Attention MLP	Linear + ReLU	200	100	ReLU
	Linear + ReLU	100	100	ReLU
	Linear	100	1	—
MLP <sub>2</sub> (Features)	Linear + ReLU	100	50	ReLU
MLP <sub>3</sub> (Value Head)	Linear + ReLU	56	150	ReLU
	Linear + ReLU	150	100	ReLU
	Linear + ReLU	100	100	ReLU
	Linear	100	1	—

Total parameters: ~78k. Inference time: ~5 ms (CPU), <1 ms (GPU).  
Input dimension 200 for attention = 100 (MLP<sub>1</sub> output) + 100 (global state).  
Input dimension 56 for MLP<sub>3</sub> = 6 (self state) + 50 (weighted features).

# VISTA-MPC with SARL



# VISTA-MPC with SARL



# VISTA-MPC with SARL

